

Machine Learning Text Classification

Hands-on introduction

Nouredine Loukil
02 May 2018



What is machine learning ?

— — —

- ML is an application of **artificial intelligence**
- ML is about automatically **learn** and improve from data
- ML do learn **without** being explicitly programmed.


Traditional programming or Machine Learning ?

— — —

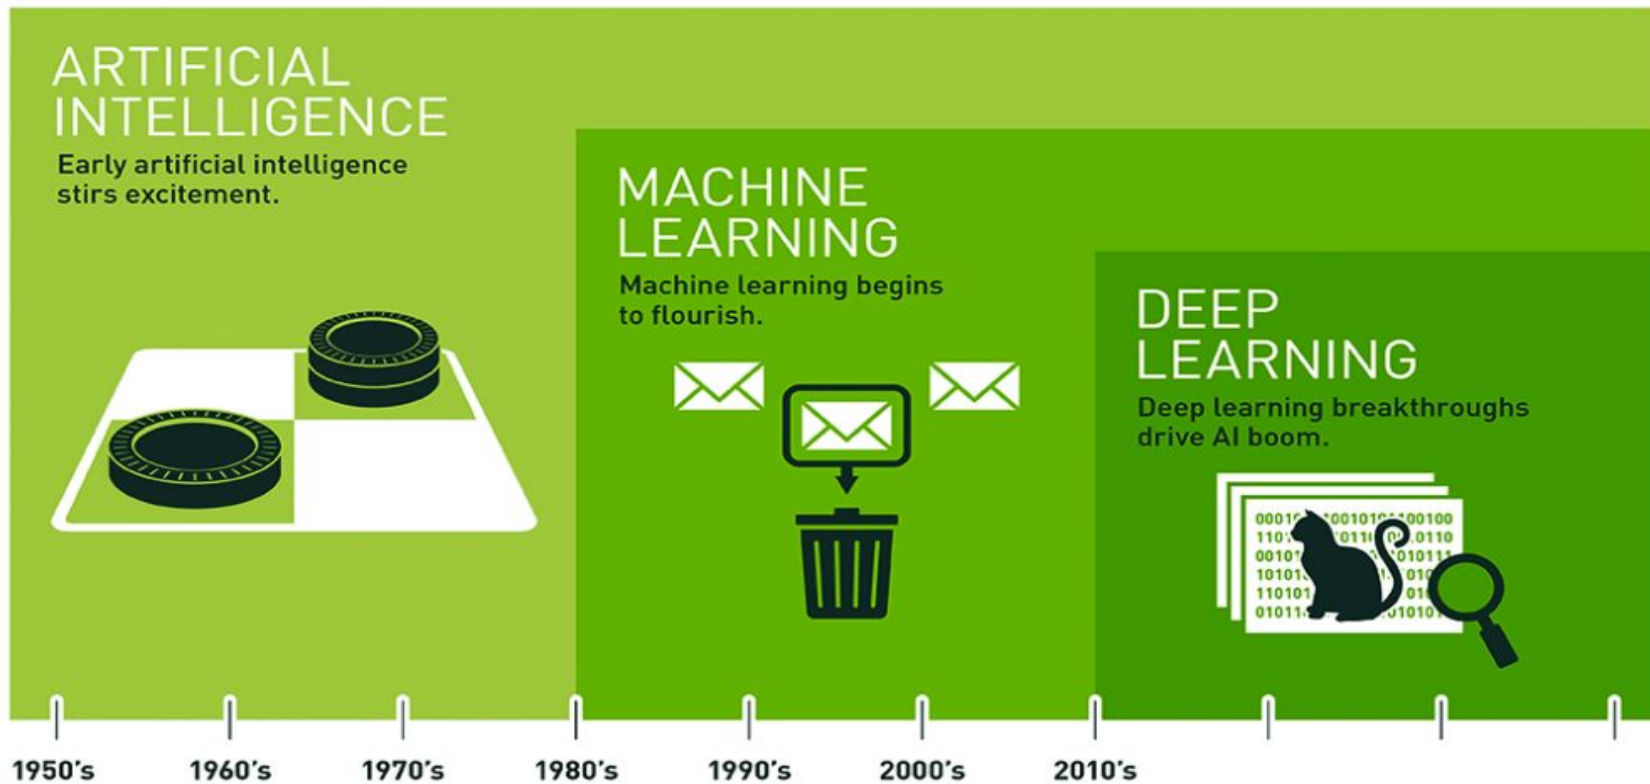
Traditional Programming



Machine Learning



History of Machine Learning

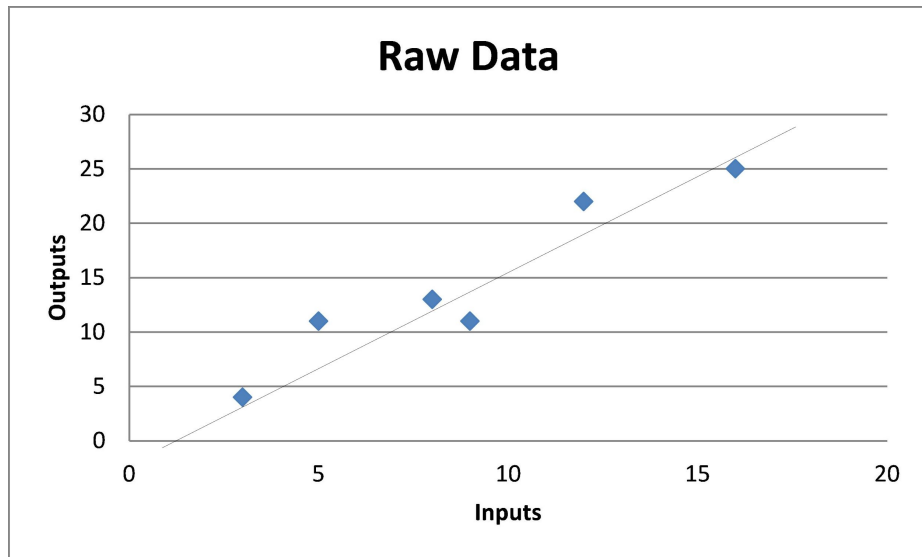


Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

Learning and prediction exercise

— — —

X Input	Y Output
3	4
5	11
8	13
9	11
13	22
17	25
10	?




$Y = aX + b$
Linear
Model

ML is trained to get a model

— — —

X Data	Y Output
3	4
5	11
8	13
9	11
13	22
17	25

Learning relations between
data (input) and output




Regression
Model

Classification of natural language text

— — —

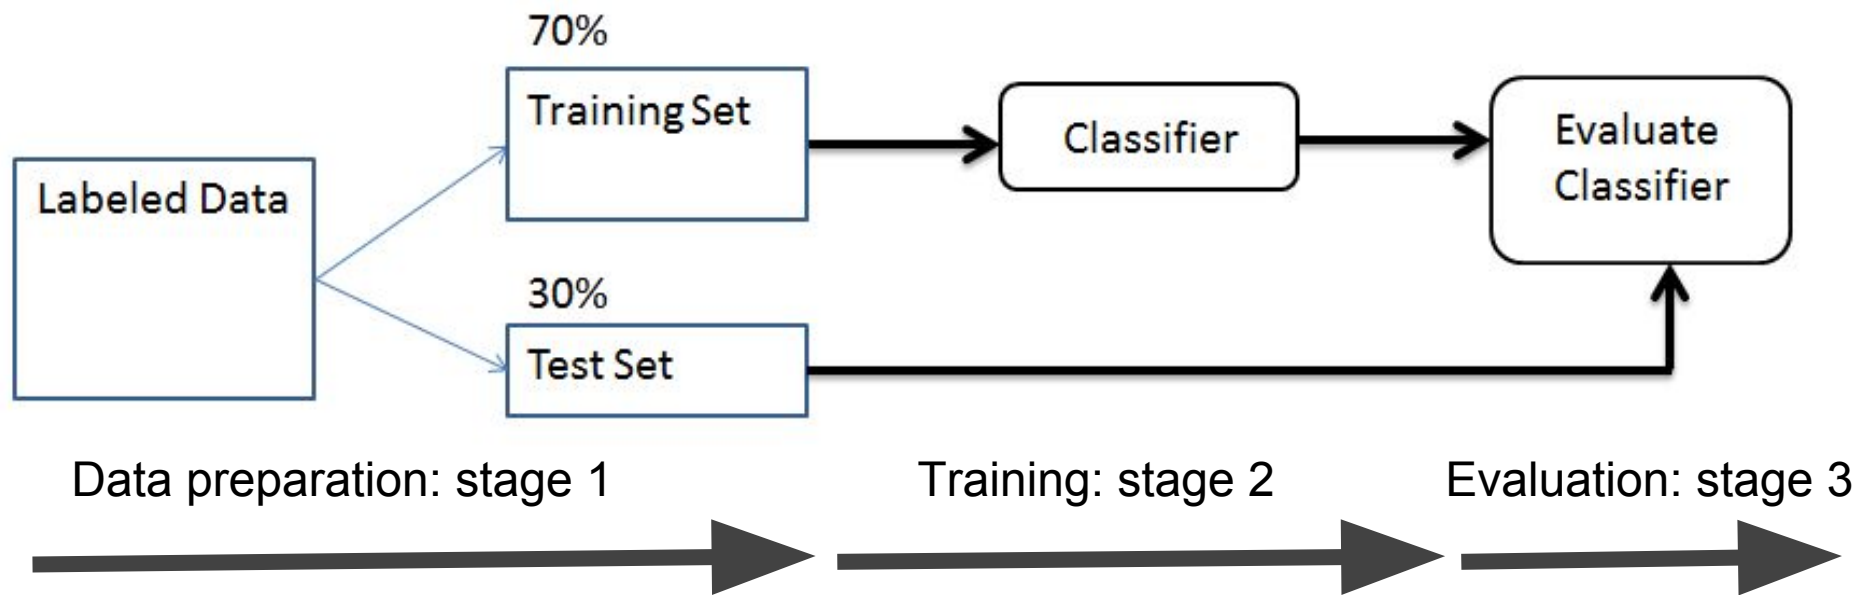
Data	Y Output
Hello my friend	+
Damn the machine	-
Oh my god	-
Thanks a lot	+
That was a shame	-
I am very angry	?

Learning relations between
data (input) and output



Classification
Model

Classification process



Classification models

— — —

1. Linear Classifiers: Logistic Regression, **Naive Bayes Classifier**
2. Support Vector Machines
3. Decision Trees
4. Boosted Trees
5. Random Forest
6. Neural Networks (deep learning)
7. Nearest Neighbor

Naive (why?) Bayes model

The diagram shows the Naive Bayes formula with four labels and arrows pointing to its components:

- Likelihood**: Points to $P(x | c)$ in the numerator.
- Class Prior Probability**: Points to $P(c)$ in the numerator.
- Posterior Probability**: Points to $P(c | x)$ on the left side of the equation.
- Predictor Prior Probability**: Points to $P(x)$ in the denominator.

$$P(c | x) = \frac{P(x | c) P(c)}{P(x)}$$

$$P(c | X) = P(x_1 | c) \times P(x_2 | c) \times \cdots \times P(x_n | c) \times P(c)$$

How Naive Bayes algorithm works?

1. Convert the data set into a frequency table
2. Create Likelihood table
3. Use Naive Bayesian equation to calculate the posterior probability for each class.
4. The class with the highest posterior probability is the outcome of prediction.

Weather	Play
Sunny	No
Overcast	Yes
Rainy	Yes
Sunny	Yes
Sunny	Yes
Overcast	Yes
Rainy	No
Rainy	No
Sunny	Yes
Rainy	Yes
Sunny	No
Overcast	Yes
Overcast	Yes
Rainy	No

Frequency Table		
Weather	No	Yes
Overcast		4
Rainy	3	2
Sunny	2	3
Grand Total	5	9

Likelihood table		
Weather	No	Yes
Overcast		4
Rainy	3	2
Sunny	2	3
All	5	9
	=5/14	=9/14
	0.36	0.64

$$P(c | x) = \frac{P(x | c) P(c)}{P(x)}$$

Likelihood $\rightarrow P(x | c)$
 Class Prior Probability $\rightarrow P(c)$
 Posterior Probability $\rightarrow P(c | x)$
 Predictor Prior Probability $\rightarrow P(x)$

$$P(c | X) = P(x_1 | c) \times P(x_2 | c) \times \dots \times P(x_n | c) \times P(c)$$

Tool suite for the data scientist

- Python vs R
- IDE (e.g. pyCharm) vs Live code (e.g. Jupyter Notebook)
- Statistical vs Neural methods
- Scikit learn library (Machine learning in python)
- Natural Language ToolKit library (NLP in python)

Scikit Learn a.k.a. sklearn

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts (classification, regression, clustering,...)
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable under BSD license
- `>>> import sklearn`
- `>>> from sklearn import something`

Natural Language ToolKit a.k.a NLTK

— — —

- leading platform for building Python programs for Natural language processing
- easy-to-use interfaces to over 50 corpora and lexical resources
- text processing libraries for classification, tokenization, stemming, tagging, parsing, semantic analysis, ...
- `>>> import nltk`
- `>>> from nltk.corpus import treebank`

Lab: News classifier

— — —

- Create an account on <https://repl.it>
- Fork or clone the Workshop material:
<https://github.com/noureddine01/MLWorkshopMaterial>
- Create MLWorkshopLab1.py on your repl account
- Train the model
- Do a prediction
- Assess the Model by yourself
- Use the model ->

Use the model

— — —

```
# new instances where we do not know the answer
```

```
Xnew, _ = [ ["new event in the weekend", "hello world"], [] ]
```

```
# make a prediction
```

```
ynew = model.predict(Xnew)
```

```
# show the inputs and predicted outputs
```

```
for i in range(len(Xnew)):
```

```
    print("X=%s, Predicted=%s" % (Xnew[i], news.target_names[ ynew[i]]))
```


Where to go from here?

- Install Jupyter Notebook <http://jupyter.org>
- Explore other data sets
http://scikit-learn.org/stable/auto_examples/index.html#dataset-examples
- Try other algorithms/models
- Try other sklearn examples
http://scikit-learn.org/stable/auto_examples/index.html#classification
- Go for real classification problems
www.kaggle.com

Thanks for your attention